Generation scaled. Discrimination didn't.

The 2026 problem isn't making AI build faster. It's making AI know if what it built is good. Here's a system that tries.

5
AGENTS THAT EVALUATE EACH OTHER

40/100
HONEST VISUAL SCORE (NOT "LOOKS GREAT!")

48h
SELF-HEALING DETECTION WINDOW

13→5
SYSTEM COMPRESSED ITSELF

EVAL LOSS

Visual quality: 8 dimensions

Playwright screenshots → Claude vision → scored evidence. Not vibes.

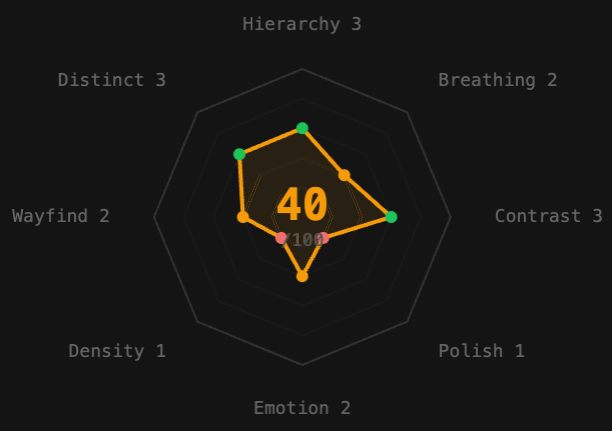

SHIP GATE

Ship readiness: NOT READY (0.35)

4 ceiling dimensions. Day-3 return is the bottleneck. This score drove the sprint plan.

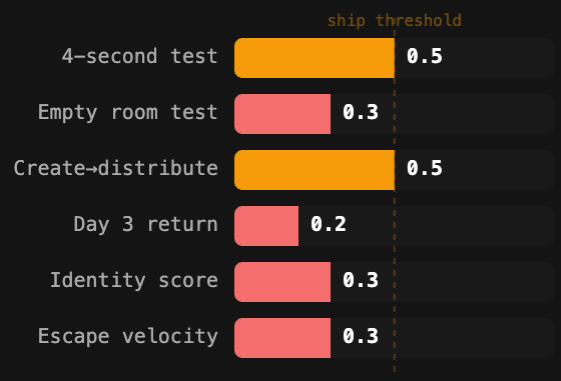

SELF-HEALING

Agent health: dead → alive

5 agents crashed silently. Meta detected, diagnosed, fixed. 2-line logs = dead. 80+ = alive.

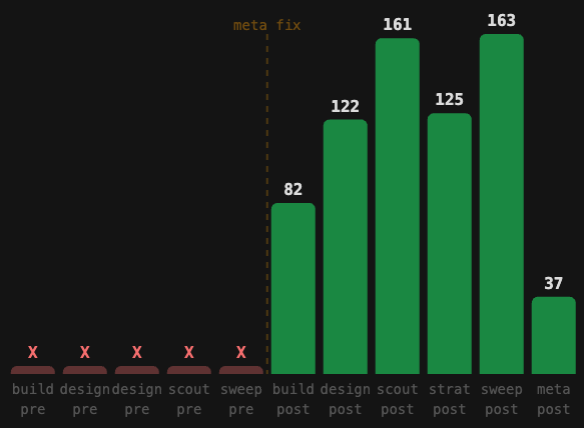

ARCHITECTURE

Self-compression: 13 → 5

The system deleted itself into shape. Less surface area, more signal.

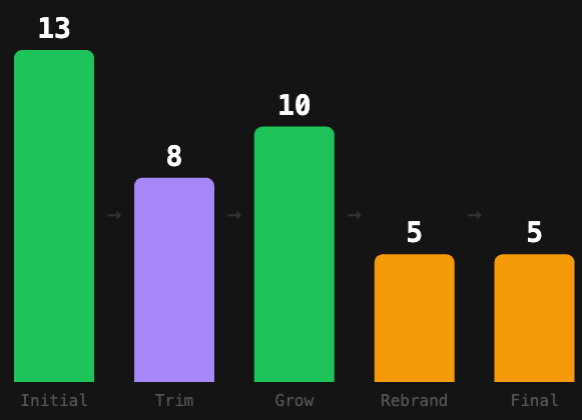

> **The pattern:** a human defines the standard in a document. The machine applies it at scale. Karpathy's program.md. Anthropic's constitution. Eval specs. Same idea. We automated the making. We haven't automated the knowing.

The loop — how evaluation drives building

Scores feed forward. Gaps become sprint priorities. The weakest dimension drives the next cycle.

Two-tier scoring

Training loss (cheap, every commit) vs eval loss (expensive, on demand). Same reason ML needs both.

	score.sh	taste.mjs
Cost	$0	~$0.50
Speed	2 sec	30 sec
Frequency	every commit	on demand
Method	grep + find	screenshots + vision
Dimensions	3	8
	training loss	eval loss

● score.sh (grep) ● taste.mjs (vision)

Meta grades agents on artifacts

Not prose quality. Did it produce its required output? Wrong calls lower credibility.

● cycle 1 ● cycle 2

Why this matters now — the convergence

Multiple independent lines of evidence point the same direction.

PATTERN

Karpathy's autoresearch

630 lines. 100 experiments overnight. The human writes program.md — the definition of "better." The agent writes code. Taste is the bottleneck, not implementation.

THEORY

Wei's Verifier's Law

AI solves every task where verification is easy. Tasks where checking is harder than producing — aesthetic judgment, product taste — are where humans keep leverage.

PROOF

Constitutional AI

Anthropic proved it: when humans define principles upstream, AI self-evaluation works. Remove the constitution, the system drifts. The document IS the product.

57% ORGS WITH AGENTS IN PRODUCTION

52% RUNNING OFFLINE EVALS

32% CITE QUALITY AS TOP BARRIER

LangChain State of Agent Engineering, 2026

The integration gap. The models are general enough. The question is whether your workflow captures what the model learned and compounds it. The gap between "I use Claude" and "I have

> agents that evaluate each other, compound knowledge, and self-correct" — that gap is entirely systems thinking, not intelligence.